

Stat 324 – Introduction to Statistics for Engineers

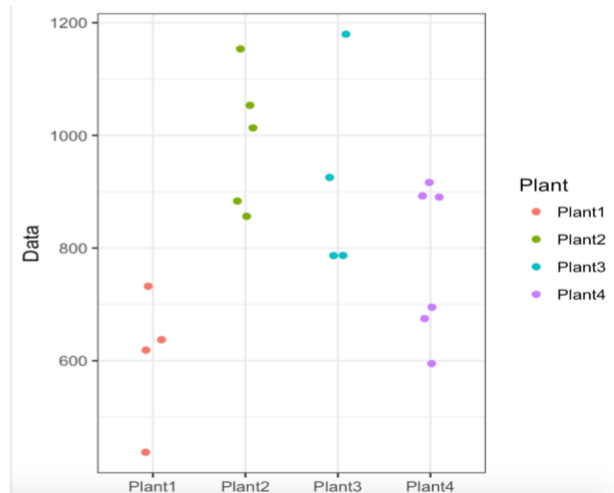
LECTURE 15: ANOVA AND KRUSKAL-WALLIS, TESTS OF CENTER FOR 2+ GROUPS
SECTIONS 8.2-8.4, 8.6, 9.3-9.5

Comparison of Multiple Groups

In the article “Review of Development and Application of CRSTER and MPTER Models” (R.Wilson *Atmospheric Environment*), several measurements of the maximum hourly concentrations of SO₂ are presented for each of four power plants. The results are as follows.

We would like to evaluate whether all four plants have similar maximum hourly concentrations of SO₂ (in the populations of SO₂s based on this sample data).

Plant 1				
438	619	732	638	
Plant 2				
857	1014	1153	883	1053
Plant 3				
925	786	1179	786	
Plant 4				
893	891	917	695	675 595



Comparison of Multiple Groups

We know how to compare pairs of populations – we could use :

independent, so need to use 2 sample t test, wilcoxon rank sum (mann-whitney U) permutation, bootstrap

The pairs (or more) of populations and their measurements are distinguished by **levels of a factor**.

The factor is a **categorical** variable. In experiments, the different levels are often different **treatments**.

E.g. The factor in the article is “Plant” and it has levels 1, 2, 3, and 4.

The experimental units on which the treatments are assigned and measured are called **replicates**.

E.g. The Plant1 level has 4 replicates, Plant2 level as 5 replicates, etc.

To compare centers across populations, we need to evaluate shifts relative to the variability we see in the replicates.

Comparison of Multiple Groups Using ANOVA

One-Way Analysis of Variance (ANOVA): When we only have 1 factor and 1 quantitative response and we are interested in whether the factor level means differ

Null: $H_0: \mu_1 = \mu_2 = \dots = \mu_T$ where μ_i is the population mean of the i 'th group. Vs

Alternative: H_A : at least one of the μ_i differ from another

Our test statistic is built by **comparing the variability**

(A) *between group means*
(~SD of group means)

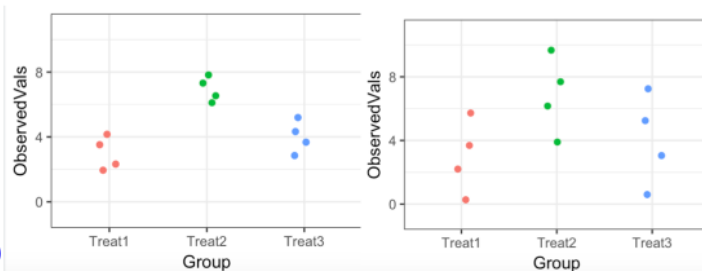
to variability

(B) *within the groups-pooled variance*
across groups

If larger variability between group means compared to small variability within groups, the stronger the evidence against the null (and therefore a larger test statistics and smaller p-value)

An example:

(same # of obs)
The treatment means are equal across both graphs, but there is more overlap between groups in the Right graph. Left graph should give us stronger evidence of difference in means (evidence null against)




One-Way Analysis of Variance (ANOVA) [nitty gritty details]

The notation we will use:

- t be the number of treatments
- i index the treatments, i from 1:4
- n_i number of replicates in treatment i , $n_1 = 4$
- $N = \sum_{i=1}^t n_i$ is total sample size, $N = 19$
- y_{ij} is the j th observation in the i th group
- $\bar{y}_{i.} = \frac{1}{n_i} \sum_{j=1}^{n_i} y_{ij}$ sample mean of the i th group, $\bar{y}_{1.} = 606.8$
 \uparrow
 average over
- $\bar{y}_{..} = \frac{1}{N} \sum_i (\sum_{j=1}^{n_i} y_{ij})$ sample mean grand mean, $\bar{y}_{..} = 827.8$

T_i						n_i	$\bar{y}_{i.}$	s_i
Plant 1	438	619	732	638		4	606.8	122.9
Plant 2	857	1014	1153	883	1053	5	992.0	122.7
Plant 3	925	786	1179	786		4	919.0	185.3
Plant 4	893	891	917	695	675	5	777.7	138.8
Overall: $N = 19, \bar{y}_{..} = 827.8$								



One-Way Analysis of Variance (ANOVA) [nitty gritty details]

Each observation's value y_{ij} can be thought of as a sum of pieces

$$\text{Observation} = \text{Grand Mean} + \text{Deviation due to Treatment} + \text{Deviation due to Obs within group (Residual)}$$

$$y_{ij} = \bar{y}_{..} + (\bar{y}_{i.} - \bar{y}_{..}) + (y_{ij} - \bar{y}_{i.})$$

Trt Group 1:


$$438 = 827.8 + (606.8 - 827.8) + (438 - 606.8)$$

$$438 = 827.8 + (-221) + (-168.8)$$

Trt Group 3:


$$1179 = 827.8 + (919 - 827.8) + (1179 - 919)$$

$$1179 = 827.8 + 91.2 + 260$$



One-Way Analysis of Variance (ANOVA) [nitty gritty details]

$$\begin{aligned}
 \text{Observation} &= \begin{bmatrix} 438 & 619 & 732 & 638 \\ 857 & 1014 & 1153 & 883 & 1053 \\ 925 & 786 & 1179 & 786 \\ 893 & 891 & 917 & 695 & 675 & 595 \end{bmatrix} \\
 \text{Grand Mean} &+ \begin{bmatrix} 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \end{bmatrix} \\
 \text{Deviation due to Treatment} &+ \begin{bmatrix} -221 & -221 & -221 & -221 \\ 164.2 & 164.2 & 164.2 & 164.2 & 164.2 \\ 91.2 & 91.2 & 91.2 & 91.2 \\ -50.1 & -50.1 & -50.1 & -50.1 & -50.1 & -50.1 \end{bmatrix} \\
 \text{Deviation due to Obs within group (Residual)} &+ \begin{bmatrix} -168.8 & 12.2 & 125.2 & 31.2 \\ -135 & 22 & 161 & -109 & 61 \\ 6 & -133 & 260 & -133 \\ 115.3 & 113.3 & 139.3 & -82.7 & -102.7 & -182.7 \end{bmatrix}
 \end{aligned}$$



$$438 = 827.8 - 221 - 168.8$$

*Slightly off due to Rounding

One-Way Analysis of Variance (ANOVA) [nitty gritty details]

$$\begin{aligned}
 (\text{Observation} - \text{Grand Mean}) &= (\text{Deviation due to Treatment} + \text{Deviation due to Obs within group}) \\
 (y_{ij} - \bar{y}_{..}) &= (\bar{y}_{i.} - \bar{y}_{..}) + (y_{ij} - \bar{y}_{i.}) \\
 \begin{bmatrix} 438 & 619 & 732 & 638 \\ 857 & 1014 & 1153 & 883 & 1053 \\ 925 & 786 & 1179 & 786 \\ 893 & 891 & 917 & 695 & 675 & 595 \end{bmatrix} - \begin{bmatrix} 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \\ 827.8 & 827.8 & 827.8 & 827.8 & 827.8 & 827.8 \end{bmatrix} &= \begin{bmatrix} -221 & -221 & -221 & -221 \\ 164.2 & 164.2 & 164.2 & 164.2 & 164.2 \\ 91.2 & 91.2 & 91.2 & 91.2 \\ -50.1 & -50.1 & -50.1 & -50.1 & -50.1 & -50.1 \end{bmatrix} + \begin{bmatrix} -168.8 & 12.2 & 125.2 & 31.2 \\ -135 & 22 & 161 & -109 & 61 \\ 6 & -133 & 260 & -133 \\ 115.3 & 113.3 & 139.3 & -82.7 & -102.7 & -182.7 \end{bmatrix}
 \end{aligned}$$

Expression holds for each cell in the arrays,

$$\begin{bmatrix} -389.8 & -208.8 & -95.8 & -189.8 \\ 29.2 & 186.2 & 325.2 & 55.2 & 225.2 \\ 97.2 & -41.8 & 351.2 & -41.8 \\ 65.2 & 63.2 & 89.2 & -132.8 & -152.8 & -232.8 \end{bmatrix} = \begin{bmatrix} -221 & -221 & -221 & -221 \\ 164.2 & 164.2 & 164.2 & 164.2 & 164.2 \\ 91.2 & 91.2 & 91.2 & 91.2 \\ -50.1 & -50.1 & -50.1 & -50.1 & -50.1 & -50.1 \end{bmatrix} + \begin{bmatrix} -168.8 & 12.2 & 125.2 & 31.2 \\ -135 & 22 & 161 & -109 & 61 \\ 6 & -133 & 260 & -133 \\ 115.3 & 113.3 & 139.3 & -82.7 & -102.7 & -182.7 \end{bmatrix}$$

And each array sums to zero! $\sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..}) = \sum_{i=1}^t \sum_{j=1}^{n_i} (\bar{y}_{i.} - \bar{y}_{..}) + \sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})$

One-Way Analysis of Variance (ANOVA) [nitty gritty details] ya, no sh:t

(Observation --- Grand Mean)² = (Deviation due to Treatment)² + (*Deviation due to Obs within group)²
(Residual²)


$$\sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..})^2 = \sum_{i=1}^t \sum_{j=1}^{n_i} (\bar{y}_{i.} - \bar{y}_{..})^2 + \sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2 =$$

$$\sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..})^2 = \sum_{i=1}^t n_i (\bar{y}_{i.} - \bar{y}_{..})^2 + \sum_{i=1}^t (n_i - 1) s_i^2 =$$

$$s_i^2 = \frac{(y_{ij} - \bar{y}_{i.})^2}{n_i - 1}$$

Total Sum of Squares = Treatment Sum of Squares (between) + Error Sum of Squares (within)

SSTOT = SSTreat + SSErr 683448.5 = 378610.4 + 304838.1 (Leaving Values Exact)




One-Way Analysis of Variance (ANOVA) [nitty gritty details]

We need to divide the squared errors by the appropriate degrees of freedom (df) so they are proper variances.

$$f_{tot} = N - 1 \quad MSTot = \frac{SSTot}{N - 1} = \frac{1}{N - 1} \sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{..})^2$$

$$f_{trt} = t - 1 \quad MSTrt = \frac{SSTrt}{t - 1} = \frac{1}{t - 1} \sum_{i=1}^t n_i (\bar{y}_{i.} - \bar{y}_{..})^2$$

$$f_{err} = N - t \quad MSE = \frac{SSErr}{N - t} = \frac{1}{N - t} \sum_{i=1}^t \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i.})^2$$



Anova Table

Source	Sum of Squares	df	MS*
Treatment	SSTrt	t-1	SSTrt/(t-1)
Error	SSErr	N-t	SSErr/(N-t)
Total	SSTot	N-1	

Anova Table

Source	Sum of Squares	df	MS*
Treatment	378610.4	4-1=3	378610.4/3
Error	304838.1	17-4=15	304838.1/15
Total	683448.5	or 17-1=16	

One-Way Analysis of Variance (ANOVA) [nitty gritty details]

Under the null, our test statistic (F) is the ratio of between variability to the within variability

$$F = \frac{MSTrt}{MSE} = \frac{SSTrt / (t-1)}{SSErr / (N-t)}$$

If $H_0: \mu_1 = \mu_2 = \dots \mu_t$ vs $H_A: \text{at least one of the } \mu_i \text{ are different from another}$

If MSTrt is large relative to MSE, then F is Bigger and we have more evidence against the null (and more evidence of a treatment effect).

If F is small, there is a lot of variability unaccounted for by the group means which suggests there is no real difference between groups.

$$F_{obs} = \frac{MSTrt_{obs}}{MSE_{obs}} = \frac{378610.4 / 3}{304838.1 / 15} = 6.210024$$

But what are sizes of the values of F that we would expect to see even if the null is true?

We need to look at F distribution to see what "extreme" is

One-Way Analysis of Variance (ANOVA) [nitty gritty details]

Under the null of our test statistic (and another big assumption we'll discuss soon...)

$$F_{(v_1=t-1, v_2=N-t)} = \frac{MSTrt}{MSE} = \frac{SSTrt / (t-1)}{SSErr / (N-t)} = \frac{\text{between treatment variability}}{\text{within treatment variability}}$$

The sizes of the values of F that we would expect to see even if the null is true is determined by the Trt and Error degrees of freedom

One-Sided Test always? yes, we will only find area above F

Looking at an F table....

$$F_{obs} = \frac{MSTrt_{obs}}{MSE_{obs}} = \frac{378610.4 / 3}{304838.1 / 15} = 6.210024$$

$$P(F_{(3,15)} > 6.21) = \text{value between } \underline{.01} \text{ and } \underline{.001}$$

Conclusion? strong evidence that not all group means are equal

Table entry for p is the critical value F^* with probability p lying to its right.




TABLE E		F critical values								
		Degrees of freedom in the numerator								
	p	1	2	3	4	5	6	7	8	9
1	.100	39.86	49.50	53.59	55.83	57.24	58.20	58.91	59.44	59.86
	.050	161.45	199.50	215.71	224.58	230.16	233.99	236.77	238.68	240.54
	.025	647.79	799.50	864.16	899.58	921.85	937.11	948.22	956.66	963.28
	.010	4052.2	4999.5	5403.4	5624.6	5763.6	5859.0	5928.4	5981.1	6022.5
2	.100	18.51	19.00	19.16	19.25	19.30	19.33	19.35	19.37	19.38
	.050	38.51	39.00	39.17	39.25	39.30	39.33	39.36	39.37	39.39
	.025	98.50	99.00	99.17	99.25	99.30	99.33	99.36	99.37	99.39
	.010	998.50	999.00	999.17	999.25	999.30	999.33	999.36	999.37	999.39

15
•.01
•.001

One-Way Analysis of Variance (ANOVA)

In R:

```
> anova<-aov(Data~Plant, data=Plantdata)
> Plantanova<-aov(Data~Plant, data=Plantdata)
> summary(Plantanova)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Plant	3	378610	126203	6.21	0.00592 **
Residuals	15	304838	20323		

(Error)

```
> Plantlm<-lm(Data~Plant, data=Plantdata)
> anova(Plantlm)
```

Analysis of Variance Table

Response: Data

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Plant	3	378610	126203	6.21	0.005917 **
Residuals	15	304838	20323		

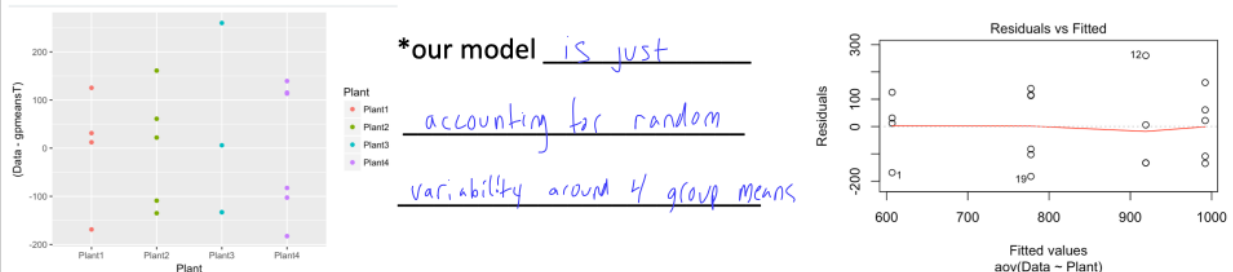
One-Way Analysis of Variance (ANOVA)

Assumptions we are making:

- (1) Groups are independent and observations within groups are independent/randomly selected
- (2) σ^2 SE is like a pooled variance between the groups, so we need to make the assumption that the variances are approximately equal in each group.

- a. Just like we did for the T-test, we can graphically evaluate this by looking at dotplots of data by group and/or
- b. Dot plots of within-group error (residuals)

Residual = Actually Observed Value - Value The *Model Predicted It Should Have Been



One-Way Analysis of Variance (ANOVA) [nitty gritty details]

Assumptions we are making:

(2) MSE is like a pooled variance between the groups, so we need to make the assumption that the variances are equal in each population.

The F test is so powerful because pooling of variances. If we believe the treatment variances are equal, we get a much more accurate estimate of the within variability if we use all the data rather than just the data from two treatments at a time, and that estimate will be consistent over all pairs. The accuracy is reflected in the increased df of an F compared to each of the possible two-sample t test.

*We can also look at ratio of SDs :

if ratio of smallest sample SD to largest sample SD is within a factor of 2 (1/2), then assumption is reasonable.


Trt1 SD: 122.9

Trt2 SD: 122.7

Trt3 SD: 185.3

Trt4 SD: 138.8

$$122.7 / 185.3 = 0.662$$



One-Way Analysis of Variance (ANOVA) [nitty gritty details]

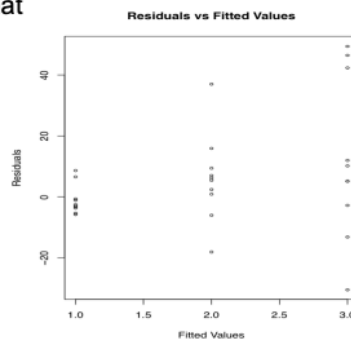
Assumptions we are making:

(2) MSE is like a pooled variance between the groups, so we need to make the assumption that the variances are equal in each populations.

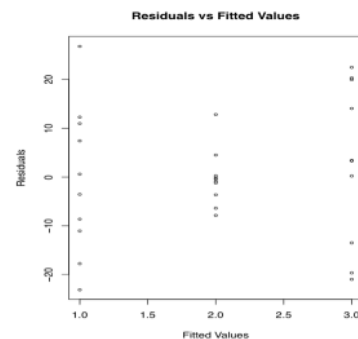
The reason we plot the residuals vs the fitted values is that it is often the case that variability will increase with increasing fitted values.

*As an extreme example, the SD of weights of 10 elephants will tend to be larger than the SD of weights of 10 hamsters. So, we like to plot against the fitted values to make it easier to see changes in variability of this type.

Here are some examples of residuals vs fitted plots that would indicate violation of constant variance:



"Funnel pattern"



Unequal variance

One-Way Analysis of Variance (ANOVA) [nitty gritty details]


Assumptions we are making:

(3) The distribution of $\frac{MSTrt}{MSE} = \frac{SSTrt/(t-1)}{SSError/(N-t)} = \frac{\text{between treatment variability}}{\text{within treatment variability}}$ is only $F_{(v_1=t-1, v_2=N-t)}$

if the population of errors (residuals) is Normal

We check this again with the QQplot (and/or histogram) of our residuals

Does this assumption seem to be well met? Not great because we have a heavy tail, we will learn rank-based next time



Multiple Comparisons Following Significant ANOVA

Conclusions we can draw:

- (1) If the F test is not significant, we do not reject the null and we're done.
- (2) If the F test is significant, we can only say at least two of the groups are significantly different (It doesn't tell us which groups though or by how much)
- (3) One way to figure these things out is with post-hoc pairwise tests that uses the pooled variance (The F test is more powerful than regular two-sample t tests because if treatment variances are actually equal, we get a much more accurate estimate of the shared variance).

To compare treatments a and b we can:

A) compare the test statistic $t = \frac{\bar{y}_a - \bar{y}_b}{\sqrt{MSE(\frac{1}{n_a} + \frac{1}{n_b})}}$ to a t_{df_E} or

B) Make a $100(1 - \alpha)\%$ CI on the difference of means: $\bar{y}_a - \bar{y}_b \pm t_{df_E, \alpha/2} \sqrt{MSE(\frac{1}{n_a} + \frac{1}{n_b})}$

*notice, you'll need to compute a different SE for the difference if there are different sample sizes in each group, but you'll get to use the t multiplier with the df_E

Multiple Comparisons Following Significant ANOVA – Fisher's Least Significant Difference (LSD)

(2) Fisher's LSD method for finding CI's for the difference in group means (after a significant ANOVA)

$df_{ERR} = 15$ and $MSE=20322.54$ and choose $\alpha = 0.05$, t multiplier = $t_{(0.025, 15)} = 2.13$

For $\mu_1 - \mu_2$: $606.8 - 992.0 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{5}\right)} = -385.2 \pm 203.831 = (-589.03, -181.369)$ *0 not included

For $\mu_1 - \mu_3$: $606.8 - 919.0 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{4}\right)} = 312.2 \pm 214.86 = (-527.06, 97.34)$ *0 not included

For $\mu_1 - \mu_4$: $606.8 - 777.7 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6}\right)} = -170.9 \pm 196.13 = (-367.04, 25.24)$ *0 included

For $\mu_2 - \mu_4$: $992.0 - 777.7 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{6}\right)} = 214.3 \pm 183.99 = (30.31, 398.29)$ *0 not included

For $\mu_2 - \mu_3$: $992.0 - 919.0 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{4}\right)} = 73 \pm 203.831 = (-130.83, 276.83)$ *0 included

For $\mu_3 - \mu_4$: $919.0 - 777.7 \pm 2.13 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6}\right)} = 141.3 \pm 196.13 = (-54.84, 337.43)$ *0 included

Anova Table

Source	Sum of Squares	df	MS*
Treatment	378610.4	4-1	378610.4 / 3 = 126203.5
Error	304838.1	19-4	304838.1 / 15 = 20322.54
Total	683448.5	19-1	

Plant 1	438	619	732	638		n_i	\bar{y}_i	s_i
Plant 2	857	1014	1153	883	1053	4	606.8	122.9
Plant 3	925	786	1179	786		5	992.0	122.7
Plant 4	893	891	917	695	675	4	919.0	185.3
					595	6	777.7	138.8
Overall: $N = 19$, $\bar{y}_\cdot = 827.8$								

ack
p'

Multiple Comparisons Following Significant ANOVA – Fisher's Least Significant Difference (LSD)

(3) We often use a lettering system to show groups that are[not] significantly different from one another (Though this make us look how different/not different they are)

Increasing means and groups not significantly different share a letter.

Treatment Mean Group (*Based on Fisher's LSD with $\alpha = 0.05$ for each CI)

Plant 1	606.8	A
Plant 4	777.7	A B
Plant 3	919.0	B C
Plant 2	992.0	C

Intervals that contained 0: $\mu_1 - \mu_4$; $\mu_2 - \mu_3$; $\mu_3 - \mu_4$

Intervals that did not contain 0: $\mu_1 - \mu_2$; $\mu_1 - \mu_3$; $\mu_2 - \mu_4$

Multiple Comparisons and Family-Wise Error Rates

Doing a large collection of tests greatly inflates our probability of making a type I error somewhere in the collection. (Think green jelly bean comic example)

- If any single test has probability α of making a Type I error (we call this the **Comparison-Wise Error Rate**, or **CW** for short), then when doing m independent tests, the probability of a Type I error on at least one test is $1 - (1 - \alpha)^m$ (we call this the **Family-Wise Error Rate**, or **FWER**).

- At a CW rate of 0.05, family-wise rate increases quickly:

m	α_{FW}
1	0.0500
2	0.0975
5	0.2262
10	0.4013

- If we want maximum power for detecting a difference and we are ok with a larger FWER than 5%, we can just leave the tests or CIs as they are.

If we want to bound our family-wise error rate, two common methods are the

a) Bonferroni Method and b) Tukey Method.

One-Way Analysis of Variance (ANOVA) – Bounding Family Error Rates

1) Bonferroni Method

This method can be used for *any* collection of m tests, not just pairwise tests following a significant ANOVA.

Simply set the individual test comparison rate to $\frac{1}{m}$ *th* of the desired family error rate, ie $\alpha_{CW} = \frac{\alpha_{FW}}{m}$

Practically, we can use this value in a few different ways:

- If using CI's, use the multiplier, $t_{df_E, (\frac{\alpha}{m})/2}$.
- If using p values, compute them as usual, but determine significance by comparing to $\frac{\alpha_{FW}}{m}$
- If using p values, multiply all p-values by m , and then determine significance by comparing to α_{FW}

CI for difference of $\mu_i - \mu_j$ is: $\bar{X}_i - \bar{X}_j \pm t_{df_E, (\frac{\alpha}{m})/2} * \sqrt{MSE \left(\frac{1}{n_i} + \frac{1}{n_j} \right)}$

In our case, the multiplier is t with $df_E = 19 - 4 = 15$ and

$$\frac{\alpha}{m} = \frac{.05}{6} * \frac{1}{2} = 0.0042 \text{ so we use the multiplier } 3.03$$

`> qt(p=.0042, df=15, lower.tail=FALSE)`
[1] 3.032375

Bonferroni is easy, but if m is very large, the power can be very small (CI very large)

One-Way Analysis of Variance (ANOVA) -Bounding Family Error Rates

(1) Bonferroni Method for finding (multiple comparison) adjusted CI's for the difference in group means

MSE=20322.54, $df_E = 19 - 4 = 15$ and $\frac{\alpha}{m} = \frac{.05}{6} * \frac{1}{2} = 0.0042$ so we use the multiplier 3.03

$$\text{For } \mu_1 - \mu_2: 606.8 - 992.0 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{5} \right)} = -385.2 \pm 289.76 = (-674.96, -95.44) * \text{no } 0$$

$$\text{For } \mu_1 - \mu_3: 606.8 - 919.0 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{4} \right)} = -312.2 \pm 305.43 = (-617.63, -6.77) * \text{no } 0$$

$$\text{For } \mu_1 - \mu_4: 606.8 - 777.7 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6} \right)} = -170.9 \pm 278.82 = (-449.72, 107.92) * \text{has } 0$$

$$\text{For } \mu_2 - \mu_4: 992.0 - 777.7 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{6} \right)} = 214.3 \pm 261.55 = (-47.25, 475.85) * \text{has } 0$$

$$\text{For } \mu_2 - \mu_3: 992.0 - 919.0 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{4} \right)} = 73 \pm 289.76 = (-216.76, 362.76) * \text{has } 0$$

$$\text{For } \mu_3 - \mu_4: 919.0 - 777.7 \pm 3.03 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6} \right)} = 141.3 \pm 278.82 = (-137.52, 420.12) * \text{has } 0$$

Treatment	Mean	Group (*Based on Bonferroni adjusted CI)
Plant 1	606.8	A
Plant 4	777.7	A B
Plant 3	919.0	B
Plant 2	992.0	B

-Notice we are seeing wider CI than unadjusted, and fewer groupings. Groups 2, 4 are no longer significantly different.

One-Way Analysis of Variance (ANOVA) -Bounding Family Error Rates

(2) Tukey Method for finding (multiple comparison) adjusted CI's for the difference in group means is a little more work, but better if we want to look at all pairwise tests (typically. more powerful than Bonferroni while still making adjustment.)

For a Tukey CI,

we use a multiplier of $\frac{Q_{\alpha_{FW}, t, df_E}}{\sqrt{2}}$ where Q is called the Studentized Range Distribution.

We will use an R function to get Q_{α_{FW}, t, df_E} for us: `> qtukey(p=(1-0.05), nmeans=4, df=(19-4)) / sqrt(2)`
[1] 2.882149

$$\text{CI for difference of } \mu_i - \mu_j \text{ is: } \bar{X}_i - \bar{X}_j \pm \frac{q_{T, N-T, \alpha}}{\sqrt{2}} * \sqrt{MSE \left(\frac{1}{n_i} + \frac{1}{n_j} \right)}$$

Reject $H_0: \mu_i - \mu_j = 0$ iff the interval doesn't contain 0 or $|\bar{X}_i - \bar{X}_j| > \frac{q_{T, N-T, \alpha}}{\sqrt{2}} * \sqrt{MSE \left(\frac{1}{n_i} + \frac{1}{n_j} \right)}$

One-Way Analysis of Variance (ANOVA) -Bounding Family Error Rates

(2) Tukey Method for finding (multiple comparison) adjusted CI's for the difference in group means

$$\begin{aligned}
 \text{For } \mu_1 - \mu_2: & 606.8 - 992.0 \pm \frac{4.075974}{\sqrt{2}} \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{5} \right)} = -385.2 \pm 274.403 = (-659.6, -110.80) * \text{no } 0 \\
 \text{For } \mu_1 - \mu_3: & 606.8 - 919.0 \pm 2.88 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{4} \right)} = -312.2 \pm 269.25 = (-601.44, -22.95) * \text{no } 0 \\
 \text{For } \mu_1 - \mu_4: & 606.8 - 777.7 \pm 2.88 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6} \right)} = -170.9 \pm 264.04 = (-434.94, 93.14) * \text{has } 0 \\
 \text{For } \mu_2 - \mu_4: & 992.0 - 777.7 \pm 2.88 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{6} \right)} = 214.3 \pm 247.696 = (-33.40, 462.0) * \text{has } 0 \\
 \text{For } \mu_2 - \mu_3: & 992.0 - 919.0 \pm 2.88 \sqrt{20322.54 \left(\frac{1}{5} + \frac{1}{4} \right)} = 73 \pm 274.4034 = (-201.40, 347.40) * \text{has } 0 \\
 \text{For } \mu_3 - \mu_4: & 919.0 - 777.7 \pm 2.88 \sqrt{20322.54 \left(\frac{1}{4} + \frac{1}{6} \right)} = 141.3 \pm 264.04 = (-122.74, 405.34) * \text{has } 0
 \end{aligned}$$

Treatment	Mean	Group (*Based on Tukey adjusted CI)
Plant 1	606.8	A
Plant 4	777.7	A B
Plant 3	919.0	B
Plant 2	992.0	B

-Notice we are seeing slightly narrower CI but same groupings as Bonferroni.

One-Way Analysis of Variance (ANOVA) -Bounding Family Error Rates

Checking our work in R:

Tukey Groupings:

Treatment	Mean	Group
Plant 1	606.8	A
Plant 4	777.7	AB
Plant 3	919.0	B
Plant 2	992.0	B

```

> TukeyHSD(Plantanova)
Tukey multiple comparisons of means
 95% family-wise confidence level

Fit: aov(formula = Data ~ Plant, data = Plantdata)

$Plant
      diff      lwr      upr    p adj
Plant2-Plant1 385.2500 109.62944 660.87056 0.0053783
Plant3-Plant1 312.2500  21.72042 602.77958 0.0332173
Plant4-Plant1 170.9167 -94.29934 436.13268 0.2866336
Plant3-Plant2 -73.00000 -348.62056 202.62056 0.8695632
Plant4-Plant2 -214.3333 -463.12800  34.46134 0.1036269
Plant4-Plant3 -141.3333 -406.54934 123.88268 0.4420141
  
```

Kruskal-Wallis (Rank-based comparison of centers)

If the residuals give strong evidence group populations aren't normal, we again use the ranks to do an overall test of difference between groups, called a Kruskal-Wallis. If this is significant, we follow it by pairwise Wilcoxon Rank Sum/Mann Whitney tests.

Assumptions:

1. n_1 observations randomly selected from population 1,
 n_2 observations randomly selected from population 2, etc

2. The populations are identical except for a possible difference in shift parameter. (have same shape and spread)

Computation of test statistic is complicated. As for Mann-Whitney, start by ranking all the data, regardless of treatment. Again assign average ranks if there are ties. The test statistic is based on how similar the average ranks are in each of the groups.

```
> kruskal.test(Data~Plant, data=Plantdata)
```

We'll use the `kruskal.test()`

Kruskal-Wallis rank sum test

Command in R.

data: Data by Plant

Kruskal-Wallis chi-squared = 9.2913, df = 3, p-value = 0.02566

Kruskal-Wallis (Rank-based comparison of centers)

Since the overall Kruskal-Wallis was significant, we have evidence that at least 1 population is shifted from another population.

Lets do Wilcoxon Rank Sum pairwise comparisons

Wilcoxon rank sum test

data: Plant1 and Plant2
W = 0, p-value = 0.01587

data: Plant1 and Plant3
W = 0, p-value = 0.0294

data: Plant1 and Plant4
W = 5, p-value = 0.1714

data: Plant2 and Plant3
W = 13, p-value = 0.5386

data: Plant2 and Plant4
W = 24, p-value = 0.1255

data: Plant3 and Plant4
W = 18, p-value = 0.2395

Notice, when we compare each to p value to 0.05, we find have evidence of a significant difference in shift between Plants 1 and 2 and between Plants 1 and 3 (similar to Bonferonni and Tukey before).

When we adjust for the multiple comparisons and use the Bonferroni adjustment (compare each p value to $0.5/6 = 0.0083$), none of the pairwise comparisons are significant.

Because Wilcoxon is only using ranks, with the small sample sizes, it has low power to find a difference between the groups.